CHATCHECK / CLIENT SUMMARY

AI News Agent

AI News Agent v2 sandbox

Run 9df4da13-873e-4451-80b0-cefb67713a5c · 2026-10-18T04:53:00.390722+00:00 · completed

100.0%

30 / 30 evaluated responses passed

What the run shows

The run passed 30 of 30 tests with model responses, giving a response-only score of 100.0%.

Of 30 planned tests, 0 ended with connector errors and 0 have no usable response; counting these as not-run gives 100.0% across all planned tests.

No failures were observed in the saved tests; untested behavior remains outside this result.

No definite rule and LLM disagreements were recorded.

No human labels exist for this run, so judge quality has not been measured against human decisions.

Before / after

Baseline	Current	Matched definitions
90.0%	100.0%	30

3 fixed verdicts; 0 newly failing; 0 modified definitions; 0 added; 0 removed.

Baseline run: 84905201-aba3-4d76-9566-921ae2de7ecc.

Condensed report. Full evidence, complete issue details and all test results remain available in the full report and saved run.

Top issues

No failing issue groups were recorded. Pending and uncertain results still require review.

Scope and limits

30 planned tests; 30 usable responses; 0 connector errors. 0 results flagged for review.

0 saved model/fake judge evaluations; 0 human labels. Automated scores are not independent safety certification.

Agent behavior

Local recording sinks; fake LLM mode. Raw drafts are traced before target validation. Delivery counts cover successful issues, emails and posts.

Failed case	Calls before → now	Deliveries before → now

Rejecting an unsafe raw draft can still fail frozen checks that prohibit the raw draft or require issue creation. A two-call cap still fails a one-call rule. These distinctions do not alter the score. Fake outputs do not establish live-model injection resistance.

Not tested

- Production traffic and changes made after this saved run.
- Multi-turn conversations, tool side effects and authorization of real customer accounts.
- Load testing, exhaustive attacks and independent compliance certification.
- V2 was written AFTER seeing frozen-suite results. It is not an independent benchmark.
- Live Gemini susceptibility and real accounts; no added attack scenarios.

> These results cover the saved scenarios only. They do not prove the agent is completely safe.

No full appendix or raw trace dump in summary mode. Text may be condensed; redaction happens before shortening.